
Notation

Inputs What are we feeding the model with. A single input is denoted with \mathbf{x} , a batch of inputs with \mathbf{X} . For example in CIFAR-10 \mathbf{x} is single image 32×32 with 3 color channel (a tensor $3 \times 32 \times 32$, channel \times height \times width). A batch of 64 images will be denoted \mathbf{X} and is a tensor $64 \times 3 \times 32 \times 32$. $|\mathbf{X}|$ is the size of batch dimension.

Classes We are concern with classification task. Let K be the set of classes and $|K|$ its cardinality. For example in CIFAR-10 $K = \{\text{airplane}, \dots, \text{truck}\} = \{k_1, \dots, k_{10}\}$ and $|K| = 10$. Let $g_{(\mathbf{x})} \in K$ be the class associated to the input \mathbf{x} (ground-truth class). When the \mathbf{x} is clear from the context the index for g will be dropped.

Labels Labels are the numerical encoding of classes. Each \mathbf{x} belong to only one class g . In this setting one popular encoding scheme is *one-hot encoding*.

$$\text{one-hot} : K \longrightarrow \{0, 1\}^{|K|} : g_{(\mathbf{x})} \longmapsto \mathbf{y}_{(\mathbf{x})} \quad \text{with components} \quad \mathbf{y}_j := \delta_{g,j}$$

Applying an encoding scheme to the batch of labels produce $\mathbf{Y}_{(\mathbf{X})}$.

Outputs What model return. Let $\mathbf{z}_{(\mathbf{x})} \in \mathbb{R}^{|K|}$ be the outputs of the model given the input \mathbf{x} . To output prediction as probability vector, usually denoted by $\hat{\mathbf{y}}$, \mathbf{z} must be normalized.

$$P_K : \mathbb{R}^{|K|} \rightarrow [0, 1]^{|K|} : \mathbf{z}_{(\mathbf{x})} \mapsto \hat{\mathbf{y}}_{(\mathbf{x})} = P_K(\mathbf{z}_{(\mathbf{x})}) \quad \text{with} \quad P_K(\mathbf{z})_j := \frac{e^{z_j}}{\sum_{k \in K} e^{z_k}}$$

$P_K(\cdot)$ is an alternative notation for the usual *softmax* function. Be more explicit about the underlying classes K give us more flexibility for constructing custom loss function.

Losses In the following losses are define as a pointwise operation that take \mathbf{x} and its corresponding label \mathbf{y} as input and return a scalar. Pointwise function apply to a batch can be combined with a reduction operator \bigoplus (e.g. sum or mean).

$$\ell : (\mathbf{x}, \mathbf{y}) \longmapsto \ell(\mathbf{x}, \mathbf{y}) \quad \mathcal{L} : (\mathbf{X}, \mathbf{Y}) \longmapsto \mathcal{L}(\mathbf{X}, \mathbf{Y}) := \bigoplus_{(\mathbf{x}, \mathbf{y}) \in (\mathbf{X}, \mathbf{Y})} \ell(\mathbf{x}, \mathbf{y})$$

Cross Entropy Loss (XE)

Standard Cross-Entropy only uses the output vector component corresponding to the ground-truth label (i.e. the g component of the output vector). Training process force $\hat{\mathbf{y}}_g$ to increase and, due to normalization, other components shrinks.

$$\ell_{\text{XE}} := -\mathbf{y}^\top \log \hat{\mathbf{y}} = -\log \hat{\mathbf{y}}_g \quad \longrightarrow \quad \mathcal{L}_{\text{XE}} := \frac{1}{|\mathbf{X}|} \sum \ell_{\text{XE}} \quad (1)$$

Complement Entropy Loss (CE)

[Che+19a] introduce *Complement Entropy Loss*. CE take as input the output vector \mathbf{z} , remove its g component and renormalise it obtaining $P_{K \setminus \{g\}}(\mathbf{z})$. CE is minus the Shannon Entropy $H(\cdot)$ of this new vector.

$$\ell_{\text{CE}} := -H(P_{K \setminus \{g\}}(\mathbf{z})) \quad \longrightarrow \quad \mathcal{L}_{\text{CE}} := \frac{1}{|\mathbf{X}|} \sum \ell_{\text{CE}} \quad (2)$$

Maximise the entropy of a distribution force it towards a flatter one (in these setting the distribution with maximum entropy is the uniform distribution). CE is paired with XE using a custom training loop:

Algorithm 1 Custom Training Loop for CE

```

1: for step in steps do
2:   Update parameters of the model using  $\mathcal{L}_{\text{XE}}$ 
3:   Update parameters of the model using  $\mathcal{L}_{\text{CE}}$ 
4: end for

```

They also proposed an empirical modification to balance the contribution that came from CE ¹.

$$\ell'_{\text{CE}} := \frac{\ell_{\text{CE}}}{|K| - 1} \quad \longrightarrow \quad \mathcal{L}'_{\text{CE}} := \frac{1}{|\mathbf{X}|} \sum \ell'_{\text{CE}}$$

Guided Complement Entropy Loss (GCE)

[Che+19b] improved upon CE by adding an additional term in the ℓ_{CE} that leverage the variation in model confidence during the training phase, i.e.

$$\ell_{\text{GCE}} := [\hat{\mathbf{y}}_g]^\alpha \ell_{\text{CE}} \quad \longrightarrow \quad \mathcal{L}_{\text{GCE}} := \frac{1}{|\mathbf{X}|} \sum \ell_{\text{GCE}} \quad (3)$$

The new factor $[\hat{\mathbf{y}}_g]^\alpha$ is called the *guiding factor*. At the beginning of training it is small (the model outputs low probability for the g -component of $\hat{\mathbf{y}}$) and then it increase with training (the model get better and the g -component of $\hat{\mathbf{y}}$ will be higher). α is a fixed hyperparameter ($\alpha = 0.2$ works reasonably well).

GCE is the only loss used in a standard training loop, i.e.

Algorithm 2 Standard Training Loop for GCE

```

1: for step in steps do
2:   Update parameters of the model using  $\mathcal{L}_{\text{GCE}}$ 
3: end for

```

Similar to [Che+19a], they modify \mathcal{L}_{GCE} to account for the number of classes

$$\ell'_{\text{GCE}} := \frac{\ell_{\text{GCE}}}{\log(|K| - 1)} \quad \longrightarrow \quad \mathcal{L}'_{\text{GCE}} := \frac{1}{|\mathbf{X}|} \sum \ell'_{\text{GCE}}$$

¹In the paper they propose a rescaling factor of $|K| - 1$ but in the source code they use $|K|$. The same inconsistency appears for $\mathcal{L}'_{\text{GCE}}$.

Hierarchical Complement Entropy (HCE)

[Che+19c] try to exploit hierarchical labels in CIFAR-100. Let G be a set that contains the siblings classes that belong to the same parental class of the ground-truth class, that is $g \in G$ and $G \subseteq K$. The *Hierarchical Complement Entropy* is

$$\ell_{\text{HCE}} := -H(P_{G \setminus \{g\}}(\mathbf{z})) - H(P_{K \setminus \{G\}}(\mathbf{z})) \longrightarrow \mathcal{L}_{\text{HCE}} := \frac{1}{|\mathbf{X}|} \sum \ell_{\text{HCE}} \quad (4)$$

The sum of HCE and XE is employed is a standard training loop:

Algorithm 3 Standard Training Loop for XE + HCE

```

1: for step in steps do
2:   Update parameters of the model using  $\mathcal{L}_{\text{XE}} + \mathcal{L}_{\text{HCE}}$ 
3: end for

```

Results

Model	Val Acc	Val Adv Acc
Che+19a_ResNet110_CIFAR10_XE	0.8698	0.5571
Che+19a_ResNet110_CIFAR10_XE,CE	0.8836	0.5923
Che+19b_ResNet56_CIFAR10_XE	0.8983	0.5333
Che+19b_ResNet56_CIFAR10_GCE	0.9005	0.6213

Table 1: CIFAR-10 experiments. Accuracy are calculate on validation dataset. FGSM based on XE is used to produce adversarial inputs.

Model	Val Acc	Val Adv Acc
Che+19a_ResNet110_CIFAR100_XE	0.6497	0.3128
Che+19a_ResNet110_CIFAR100_XE,CE	0.6516	0.3214
Che+19b_ResNet56_CIFAR100_XE	0.6001	0.2744
Che+19b_ResNet56_CIFAR100_GCE	0.6246	0.3674
Che+19c_ResNet56_CIFAR100_XE,CE	0.6164	0.2674
Che+19c_ResNet56_CIFAR100_XE,HCE	0.6164	0.3013
Ours_ResNet56_CIFAR100_HGCE	0.6356	0.3603

Table 2: CIFAR-100 experiments. Accuracy are calculate on validation dataset. FGSM based on XE is used to produce adversarial inputs.

References

- [Che+19a] Hao-Yun Chen et al. “Complement Objective Training”. In: (Mar. 2019). eprint: 1903.01182v2. URL: <http://arxiv.org/abs/1903.01182v2>.
- [Che+19b] Hao-Yun Chen et al. “Improving Adversarial Robustness via Guided Complement Entropy”. In: (Mar. 2019). eprint: 1903.09799v3. URL: <http://arxiv.org/abs/1903.09799v3>.
- [Che+19c] Hao-Yun Chen et al. “Learning with Hierarchical Complement Objective”. In: (Nov. 2019). eprint: 1911.07257v1. URL: <http://arxiv.org/abs/1911.07257v1>.